

```
In [ ]: # install dependencies
%pip install -q \
    matplotlib \
    pandas \
    pycaret \
    'pycaret[analysis]' \
    seaborn
```

WARNING: visions 0.7.5 does not provide the extra 'type-image-path'

Note: you may need to restart the kernel to use updated packages.

```
In [ ]: # global parameters
DATA_DIR = '../datasets/swell/final'
TEST_DATA_NAME = 'test-custom-1'
DO_SAVE_RESULTS = True
DO_COMPARE_MODELS = True
DO_PLOT_DATA = True
```

```
In [ ]: # set up the environment
import os
os.environ['PYCARET_CUSTOM_LOGGING_LEVEL'] = 'CRITICAL'

import pandas as pd
pd.set_option('display.max_columns', 128)
```

```
In [ ]: # prepare the data
from pathlib import Path
from pycaret.datasets import get_data
from zipfile import ZipFile

DATA = {
    name: None
    for name in ['train', TEST_DATA_NAME]
}

for data_name in DATA.keys():
    data_path = Path(DATA_DIR).joinpath(data_name)
    # extract the compressed data files
    ZipFile(data_path.with_suffix('.zip'), 'r').extract(
        str(data_path.with_suffix('.csv')), '..'
    )
    print(f'Data file "{data_name}" has been extracted successfully')
    # load the data
    print(f'Loading data file "{data_name}"')
    DATA[data_name] = get_data(dataset=f'{data_path}')
```

Data file "train" has been extracted successfully

Loading data file "train"

	MEAN_RR	MEDIAN_RR	SDRR	RMSSD	SDSD	SDRR_RMSSD	HR	pNN1
0	885.157845	853.763730	140.972741	15.554505	15.553371	9.063146	69.499952	11.1333
1	939.425371	948.357865	81.317742	12.964439	12.964195	6.272369	64.363150	5.6000
2	898.186047	907.006860	84.497236	16.305279	16.305274	5.182201	67.450066	13.0666
3	881.757865	893.460030	90.370537	15.720468	15.720068	5.748591	68.809562	11.8000
4	809.625331	811.184865	62.766242	19.213819	19.213657	3.266724	74.565728	20.2000

Data file "test-custom-1" has been extracted successfully

Loading data file "test-custom-1"

	MEAN_RR	MEDIAN_RR	SDRR	RMSSD	SDSD	SDRR_RMSSD	HR	pNN25
0	877.908621	877.615570	88.497648	13.151277	13.149960	6.729206	68.344243	6.422018
1	880.146934	879.724925	87.158728	13.135020	13.133364	6.635599	68.170436	6.422018
2	881.829910	881.458160	85.334954	13.187250	13.186095	6.471020	68.040332	6.422018
3	883.212753	884.504020	84.331627	13.186513	13.185818	6.395294	67.933802	6.422018
4	884.964151	885.190255	83.194489	13.303572	13.303106	6.253545	67.799357	6.672227

```
In [ ]: # set column specifications
TARGET_NAME = 'condition'
IGNORE_NAMES = ['datasetId']
```

```
In [ ]: # establish an experiment
from pycaret.classification import ClassificationExperiment

exp = ClassificationExperiment()
exp.setup(
    data=DATA['train'],
    test_data=DATA[TEST_DATA_NAME],
    target=TARGET_NAME,
    ignore_features=IGNORE_NAMES,
    index=False,
    session_id=123,
    remove_multicollinearity=True,
    multicollinearity_threshold=0.999,
    imputation_type=None,
)
exp.dataset_transformed.head(5)
```

	Description	Value
0	Session id	123
1	Target	condition
2	Target type	Multiclass
3	Target mapping	interruption: 0, no stress: 1, time pressure: 2
4	Original data shape	(408550, 36)
5	Transformed data shape	(408550, 29)
6	Transformed train set shape	(369289, 29)
7	Transformed test set shape	(39261, 29)
8	Ignore features	1
9	Numeric features	34
10	Preprocess	True
11	Imputation type	None
12	Remove multicollinearity	True
13	Multicollinearity threshold	0.999000
14	Fold Generator	StratifiedKFold
15	Fold Number	10
16	CPU Jobs	-1
17	Use GPU	False
18	Log Experiment	False
19	Experiment Name	clf-default-name
20	USI	f795

Out[]:

	MEAN_RR	MEDIAN_RR	RMSSD	SDRR_RMSSD	HR	pNN25	pNN50	
0	885.157837	853.763733	15.554504	9.063146	69.499954	11.133333	0.533333	199.061
1	939.425354	948.357849	12.964439	6.272368	64.363152	5.600000	0.000000	114.634
2	898.186035	907.006836	16.305279	5.182201	67.450066	13.066667	0.200000	118.939
3	881.757874	893.460022	15.720469	5.748590	68.809563	11.800000	0.133333	127.318
4	809.625305	811.184875	19.213820	3.266724	74.565727	20.200001	0.200000	87.718

5 rows × 29 columns

In []:

```
('Removed columns: ', list(
    set(IGNORE_NAMES) |
    (set(exp.dataset.columns) - set(exp.dataset_transformed.columns))
))
```

Out[]:

```
('Removed columns: ',
 ['datasetId',
  'SDRR',
  'KURT_REL_RR',
  'SD1',
  'SKEW_REL_RR',
  'SDSD',
  'RMSSD_REL_RR'])
```

```
In [ ]: # show the distributions of the data
```

```
# DO_PLOT_DATA = True
if DO_PLOT_DATA:
    # set plot parameters
    from pathlib import Path
    import matplotlib.pyplot as plt
    import seaborn as sns

    # reset old global plot parameters
    plt.rcParams()

    # adjustable global plot parameters
    COLORMAP = sns.color_palette()
    DPI = 400
    OUTLINE_WIDTH = 0.2
    plt.rcParams['axes.grid'] = False
    plt.rcParams['axes.linewidth'] = OUTLINE_WIDTH
    plt.rcParams['figure.dpi'] = DPI
    plt.rcParams['font.size'] = 4
    plt.rcParams['xtick.major.width'] = OUTLINE_WIDTH
    plt.rcParams['xtick.minor.width'] = OUTLINE_WIDTH
    plt.rcParams['ytick.major.width'] = OUTLINE_WIDTH
    plt.rcParams['ytick.minor.width'] = OUTLINE_WIDTH
    plot_dir = Path(f'../images/{TEST_DATA_NAME}')
    plot_dir.mkdir(parents=True, exist_ok=True)

    from math import ceil

    # adjustable local plot parameters
    TITLE = 'Density for transformed data'
    SUBPLOT_SIZE = (750, 300)

    # setting local plot parameters
    plots_per_col = 5
    shape = (plots_per_col, ceil(exp.dataset_transformed.shape[1] / plots_per_col))
    figsize = tuple(pxs[0] * pxs[1] / DPI for pxs in zip(SUBPLOT_SIZE, shape))
    title_params = {
        'label': TITLE,
        'fontdict': {
            'fontsize': plt.rcParams['font.size'] * 2,
            'fontweight': 'bold',
        },
        'loc': 'left',
        'pad': plt.rcParams['font.size'] * 2,
    }

    # plot grid
    axs = plt.subplots(
        nrows=shape[1],
        ncols=shape[0],
        layout='constrained',
        figsize=figsize,
    )[1].flat

    # plot target distribution
    target_dist_data = exp.y.value_counts(normalize=True)
    ax = sns.barplot(
        x=target_dist_data.index,
        y=target_dist_data.values,
        ax=axs[0],
        palette=COLORMAP,
    )

    # plot data title
```

```

axs[0].set_title(**title_params)

# plot features distribution
for x, ax in zip(exp.X_transformed.columns, axs[1:]):
    sns.histplot(
        data=exp.dataset_transformed,
        x=x,
        ax=ax,
        hue=TARGET_NAME,
        legend=False,
        linewidth=0,
        multiple='fill',
        palette=COLORMAP,
        stat='density',
    )

# save the plot
plt.savefig(
    fname=plot_dir.joinpath(f'{TITLE}.png'),
    bbox_inches='tight',
)
plt.show()

# check correlation between target and features
# adjustable plot parameters
TITLE = 'Correlations to target for transformed data'
PLOT_SIZE = (2560, 1440)

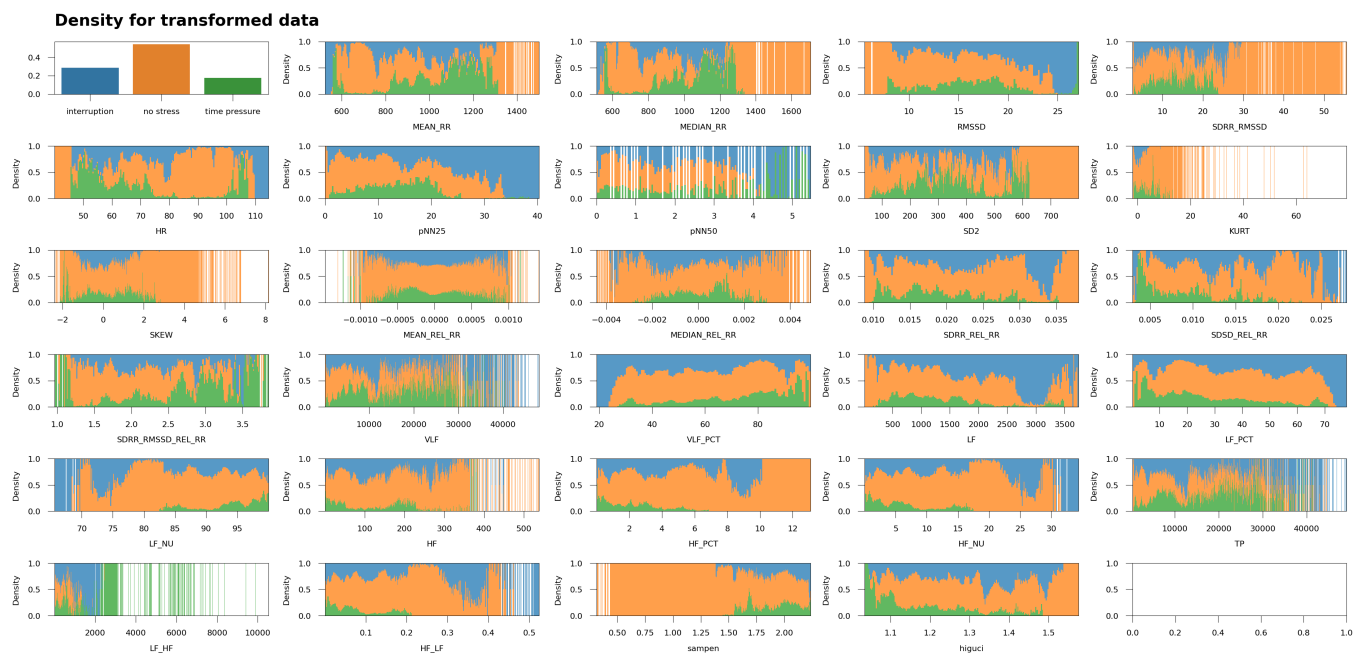
# setting plot parameters
figsize = tuple(px / DPI for px in PLOT_SIZE)
title_params = {
    'label': TITLE,
    'fontdict': {
        'fontsize': plt.rcParams['font.size'] * 2,
        'fontweight': 'bold',
    },
    'loc': 'left',
    'pad': plt.rcParams['font.size'] * 2,
}

# plot correlation to target
target_corr_data = (exp.dataset_transformed
    .corr()[TARGET_NAME]
    .drop(TARGET_NAME)
    .sort_values())
ax = target_corr_data.plot.barh(figsize=figsize)
ax.set_title(**title_params)


# save the plot
plt.savefig(
    fname=plot_dir.joinpath(f'{TITLE}.png'),
    bbox_inches='tight',
)
plt.show()

# reset plot parameters
plt.rcParamsdefaults()

```



Correlations to target for transformed data



```
In [ ]: # compare models with AUROC

# DO_COMPARE_MODELS = True
if DO_COMPARE_MODELS:
    exp.compare_models(
        exclude=[ # excludes slow and unsuitable models
            'ada',
            'catboost',
            'gbc',
            'knn',
            'lr',
            'ridge',
            'rf',
            'svm',
        ],
        sort='auc',
        cross_validation=False,
    )
None
```

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
xgboost	Extreme Gradient Boosting	0.7595	0.8793	0.7595	0.7594	0.7583	0.5917	0.5923	4.0600
et	Extra Trees Classifier	0.7249	0.8751	0.7249	0.7240	0.7241	0.5357	0.5359	4.7400
lightgbm	Light Gradient Boosting Machine	0.7293	0.8682	0.7293	0.7390	0.7301	0.5486	0.5517	5.8900
dt	Decision Tree Classifier	0.6317	0.7024	0.6317	0.6429	0.6353	0.3991	0.4006	15.2300
lda	Linear Discriminant Analysis	0.3645	0.6703	0.3645	0.5811	0.3060	0.1106	0.1695	2.2700
nb	Naive Bayes	0.4688	0.6614	0.4688	0.5309	0.4775	0.1931	0.2048	1.3000
qda	Quadratic Discriminant Analysis	0.5360	0.6395	0.5360	0.5306	0.5329	0.2094	0.2096	1.7600
dummy	Dummy Classifier	0.5444	0.0000	0.5444	0.2964	0.3838	0.0000	0.0000	1.2200

```
In [ ]: # assign the best model id manually
BEST_MODEL_ID = 'xgboost'


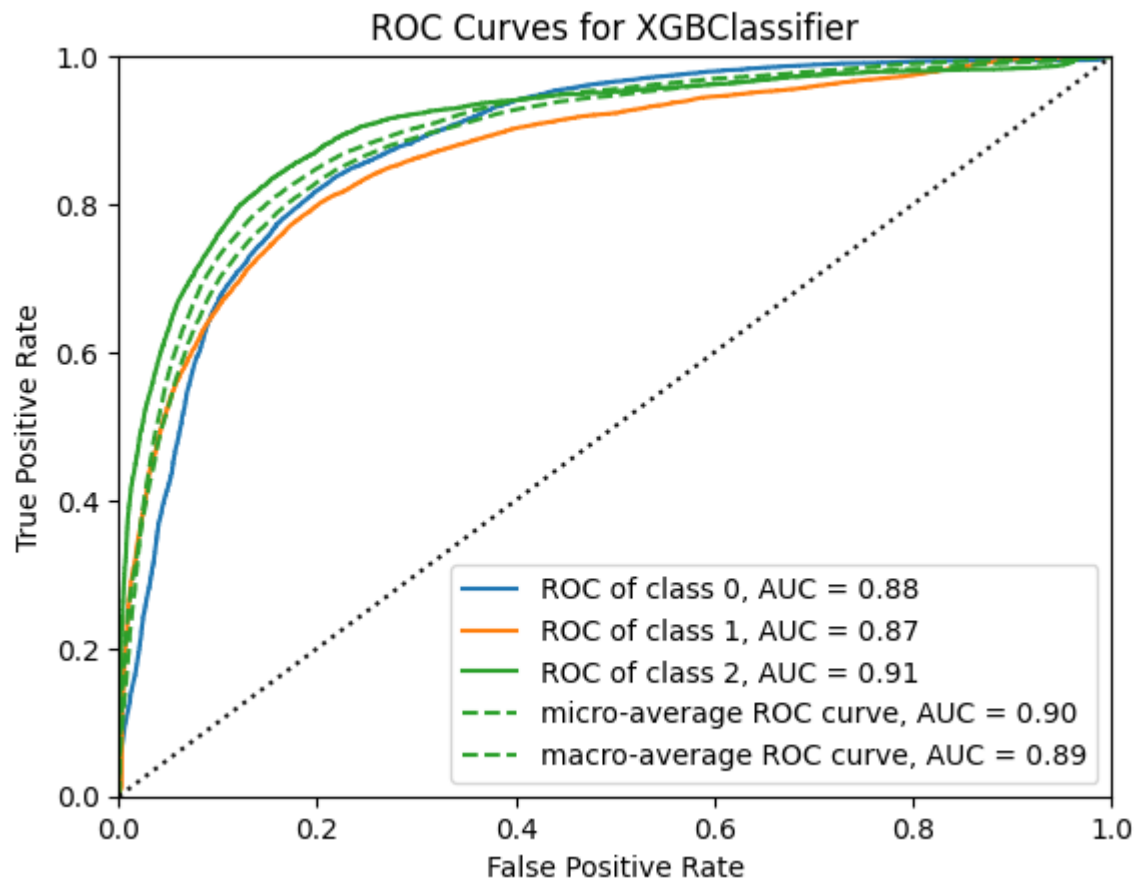
best_model = exp.create_model(
    estimator=BEST_MODEL_ID,
    cross_validation=False,
)
best_model
```

	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
Test	0.7595	0.8793	0.7595	0.7594	0.7583	0.5917	0.5923


```
Out [ ]: XGBClassifier
XGBClassifier(base_score=None, booster='gbtree', callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device='cpu', early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
```

```
In [ ]: plot_types = [
    'pipeline',
    'auc',
    'class_report',
    'feature_all',
    'parameter',
]
for plot_type in plot_types:
    exp.plot_model(
        estimator=best_model,
```

```
plot=plot_type,  
)
```



Feature Importance Plot



Parameters	
objective	multi:softprob
base_score	None
booster	gbtree
callbacks	None
colsample_bylevel	None
colsample_bynode	None
colsample_bytree	None
device	cpu
early_stopping_rounds	None
enable_categorical	False
eval_metric	None
feature_types	None
gamma	None
grow_policy	None
importance_type	None
interaction_constraints	None
learning_rate	None
max_bin	None
max_cat_threshold	None
max_cat_to_onehot	None
max_delta_step	None
max_depth	None
max_leaves	None
min_child_weight	None
missing	nan
monotone_constraints	None
multi_strategy	None
n_estimators	None
n_jobs	-1
num_parallel_tree	None
random_state	123
reg_alpha	None
reg_lambda	None
sampling_method	None
scale_pos_weight	None
subsample	None
tree_method	auto

Parameters

validate_parameters	None
---------------------	------

verbosity	0
-----------	---

```
In [ ]: # show hold-out predictions
predictions = exp.predict_model(
    estimator=best_model,
    raw_score=True,
)
display(predictions[filter(
    lambda name: name.startswith('prediction_'),
    predictions.columns,
)].sample(
    n=15,
    random_state=123,
))
predictions = None
```

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
0	Extreme Gradient Boosting	0.7595	0.8793	0.7595	0.7594	0.7583	0.5917	0.5923

	prediction_label	prediction_score_interruption	prediction_score_no stress	prediction_score_tim pressur
373100	no stress		0.0003	0.9997
373377	no stress		0.0002	0.9997
391482	interruption		0.5755	0.4221
397656	no stress		0.0000	1.0000
376683	no stress		0.0000	1.0000
373069	no stress		0.0007	0.9993
373328	no stress		0.0006	0.9994
369773	no stress		0.0000	0.9999
387397	interruption		0.9951	0.0049
386482	no stress		0.0144	0.9856
371284	no stress		0.0011	0.9259
397154	no stress		0.0000	1.0000
378792	no stress		0.0006	0.9972
398145	interruption		0.9332	0.0666
394468	no stress		0.0035	0.9964

```
In [ ]: # save the experiment and model

# DO_SAVE_RESULTS = True
if DO_SAVE_RESULTS:
    from pathlib import Path

    result_dir = Path(f'../models/{TEST_DATA_NAME}')
    result_dir.mkdir(
        parents=True,
        exist_ok=True,
    )
    exp.save_experiment(
```

```
    path_or_file=result_dir.joinpath('experiment.pkl'),  
    )  
exp.save_model(  
    model=best_model,  
    model_name=result_dir.joinpath('model'),  
    )
```

Transformation Pipeline and Model Successfully Saved